

Programming Project 1: (100 points):

Issued Date: Thu Sep 27th 2018,

Due Date: Fri Oct 12th 2018, 11.00 pm.

To be conducted in groups of 4. When you run your programs, you may compare them against other nodes in your group, if the specifications of other machines (nodes) are different. Again, please produce an extensive report that displays the results of your experimentation, describes your code, and answers all questions addressed.

Problem 1: Parallel Fractal Generation -- Investigation (15 points)

Question: The code on Mandelbrot provided (included in a separate file) computes the Mandelbrot fractal image, achieving great speedups compared to a sequential execution, by utilizing both the CPU's cores and the SIMD execution units within each core. ISPC language constructs describe *independent computations*. These computations may be executed in parallel without violating program correctness. In the case of the Mandelbrot image, computing the value of each pixel is an independent computation. With this information, the ISPC compiler and runtime system take on the responsibility of generating a program that utilizes the CPU's collection of parallel execution resources as efficiently as possible.

Program 1 - Part 1. A Few ISPC Basics (9 points)

In contrast to C, multiple program instances of an ISPC program are always executed in parallel on the CPU's SIMD execution units. The number of program instances executed simultaneously is determined by the compiler (and chosen specifically for the underlying machine). This number of concurrent instances is available to the ISPC programmer via the built-in variable `programCount`. ISPC code can reference its own program instance identifier via the built-in `programIndex`. Thus, a call from C code to an ISPC function can be thought of as spawning a group of concurrent ISPC program instances (referred to as gang). The gang of instances runs to completion, then control returns back to the calling C code.

Please familiarize yourself with ISPC language constructs by reading through the ISPC walkthrough available at <http://ispc.github.com/example.html>. The example program in the walkthrough is close to the implementation of `mandelbrot_ispc()` in `mandelbrot.ispc`.

Part 1 deliverable:

Carefully inspect the code for mandelbrot ispc (without tasks) quoted above and write a 1-2 page summary on exactly how it operates, with details for each function. Then, by connecting the problem to your group's allocated ORBIT Cores specifications **answer the following questions**: What is the maximum speedup you expect in theory (and ideally) given the specifications of your allocated CPUs? Why might the number you observe be less than this ideal (take our word for this or ... do the extra bonus questions and make the proper modifications to compile it and run it yourself)? Consider the characteristics of the computation you are performing. Describe the parts of the image that present challenges for SIMD execution. Why is this so? What adverse phenomenon/characteristic do they add to the execution (what is the name of this phenomenon/characteristic)?

Note: for such piece of code, the ISPC compiler maps gangs of program instances to SIMD instructions executed on a single core.

Extra Bonus! (TBD points): Compile and run the program mandelbrot ispc after you make the proper changes to set it up and adjust it to your environment. **CAUTION:** This may take some time and it is for those of you that are curious, hands-on, and eager enough to attempt!

Program 1 - Part 2. ISPC Tasks (6 points)

ISPCs SPMD execution model and mechanisms like `foreach` facilitate the creation of programs that utilize SIMD processing. The language also provides an additional mechanism utilizing multiple cores in an ISPC computation. This mechanism is launching *ISPC tasks*.

The `launch[2]` command in the function `mandelbrot_ispc_withtasks` launches two tasks. Each task defines a computation executed by a gang of ISPC program instances. As given by the function `mandelbrot_ispc_task`, each task computes a region of the final image. Similar to how the `foreach` construct defines loop iterations that can be carried out in any order (and in parallel by ISPC program instances, the tasks created by this launch operation can be processed in any order (and in parallel on different CPU cores).

Part 2 deliverable:

Study how `mandelbrot_ispc` can be run with the parameter `--tasks`. Do your research, do your investigation on the ispc walkthrough available and also on extra related resources and come up with an approximation for the speedup using tasks over the version of `mandelbrot_ispc` that does not partition that computation into tasks. Does the speedup depend on the parameter `tasks` and in what way?

The performance of `mandelbrot_ispc --tasks` can be improved by changing the number of tasks the code creates. By only changing code in the function `mandelbrot_ispc_withtasks()`, you should be able to achieve performance that exceeds the sequential version of the code by a great deal! How do you determine in theory how many tasks to create? Why do you think that the number you chose work best? Provide justification by theory or other counter examples.

Extra Bonus! (TBD points): Compile and run the program `mandelbrot_ispc` after you make the proper changes to adjust it to your environment with the parameter `--tasks`. What speedup do you observe? What is the speedup over the version of `mandelbrot_ispc` that does not partition that computation into tasks? By only changing code in function `mandelbrot_ispc_withtasks()`, you should be able to achieve performance that exceeds the sequential version of the code by a great deal! By how exactly for your own ORBIT cores set-up? How do you determine how many tasks to create? **CAUTION:** This may take some time and it is for those of you that are curious, hands-on, and eager enough to attempt despite the time limitations!

Problem 2: Parallel Square Root Computation (35 points)

Question: Write an ISPC program that computes the square root of 20 million random numbers (real, no necessarily integers) between 0 and 8 (*Hint:* Use Newton's method to solve the equation: $\frac{1}{x^2} - S = 0$). Please use multiple iterations for `sqrt` to converge to an accurate solution (accurate solution is one that gives $< 10^{-4}$ difference compared to the true value).

ISPCs SPMD execution model (and libraries), like “foreach”, implement SIMD processing. ISPC also provides a mechanism that “launches ISPC tasks”, utilizing this way multiple cores in an ISPC computation. Check the `launch[num]` command in function `mandelbrot_ispc_withtasks` and in the Intel ISPC Spec. This command launches *num* tasks. Each task defines a computation that will be executed by a “thread” of ISPC gangs/instances. As given by the function `mandelbrot_ispc_task`, each task computes a region of the final image in any order (in parallel on different CPU cores).

What is the ISPC implementation speedup for single CPU core (no tasks launched) and when using multiple cores (with tasks)? What is the speedup due to SIMD parallelization? What is the speedup due to multi-core parallelization? Extend your code to utilize 2, 3, 4, 5, 6, 7, 8 threads, partitioning the computation accordingly. In your write-up, produce a graph of speedup compared to a sequential implementation (which you need to implement as well) as a function of the number of cores and threads used. Is speedup linear in the number of cores used? In the number of threads used? Please justify why yes or why not.

Then write another version of the `sqrt` function using AVX intrinsics. You may consult the Intel Intrinsics Guide: <https://software.intel.com/sites/landingpage/IntrinsicsGuide/>.

If you wish to tell the ISPC compiler to generate AVX, rather than SSE instructions, please see the Makefile comment about ISPC's `--target=avx2-i32x8` compiler flag. For slightly better performance use `--target=avx2-i32x16`, which will use a gang size of 16 instances and use two AVX instructions to implement an operation for the entire gang. (more details here: <https://ispc.github.io/ispc.html#selecting-the-compilation-target>).

Problem 3: Password Cracker (50 points)

Question: Implement a simple **password cracker** using ISPC. You are going to make a password which is between 1-8 ascii characters. You will then use one of the popular hash functions to convert your password to a hash. You are expected to do adequate research on this topic and list your results in a detailed report. So, let's assume that the attacker (you-the group-the programmer) is exposed only to the hash of your password and to no other information. Another question to answer: what is a hash function? What are the properties of a hash function? When the attacker gets a hashed password can he directly know what the hash function is? Or can he retrieve or investigate the hash? How does the attacker eventually obtain the hash so that the attacker can brute-force or even use a rainbow table to try any combination of characters through the hash in order to obtain the same hash value? So, you are asked to do exactly the same. You are to implement your own password cracker based on the rules above. You are provided the range of the password length, which is to be between 1-8 characters. Provided this information, your group should brute force the password, i.e., take every possible combination of words between the mentioned length, convert it into a hash and compare it with the provided hash. Yet a couple of extra questions to answer: i) When to use a Dictionary attack vs. a Rainbow Table attack? ii) What are the resource requirements for each type of attack? Which uses more storage? Which requires more pre-computation? Which requires more analysis time? (per-hash vs. batch cracking)

This task can be easily parallelized. You are expected to use and exploit exactly the concepts worked upon in Project 1 Part 1 in the best possible way to provide the fastest times for your password cracker. And you may select various and diverse ways of your problem decomposition and assignment with the end goal to achieve optimal workload utilization which may of course – under circumstances – which do you believe? – lead also to fastest times of password cracking. To facilitate your understanding on the concept of assignment, let's consider a scenario for which thread 1 checks only combinations of length 1, thread 2 does that for 2 letter words, thread 3 for 3 letter words etc. So, you first generate the hashes and then matching them in bulk via your dedicated hardware (what type of HW in particular should you be using here?).

Similar to Part 1: What is the ISPC implementation speedup for single CPU core (no tasks launched) and when using multiple cores (with tasks)? What is the speedup due to SIMD

parallelization? What is the speedup due to multi-core parallelization? Extend your code to utilize 2, 3, 4, 5, 6, 7, 8 threads, partitioning the computation accordingly. In your write-up, produce a graph of speedup compared to a sequential implementation (which you need to implement as well) as a function of the number of cores and threads used. Is speedup linear in the number of cores used? In the number of threads used? Please justify why yes or why not.

You may also conduct some search on the web and utilize in addition an existing software tool for password cracking. Does your tool lend itself easily to parallel execution? You are expected to do some research and find one that does. Conduct a similar experiment as the one with your own password cracker. Compare the results and discuss them in your report.

Note: Also experiment with a number of different password characters and report the difference in the timing results.

All groups in addition to your own imaginary passwords, also use the following to compare against all groups and all varying implementations: bv37qi#f. Apply a number of popular hashes and obtain a number of versions of the password. And you take it from there...

You are encouraged to run your code on the orbit testbed in WINLAB. Use your ORBIT account per group. Time slots are reserved everyday normally between 20.00-12.00am but also randomly earlier in the day. The ISPC compiler is already installed in “mariasfirstimage” but you may download ISPC at <http://ispc.github.com/>.